

Computer Vision for HCI

Image Segmentation and Template Matching

1

Image Segmentation

- Goal: Partition an image into distinct regions containing each pixels with similar attributes
- Use “*discontinuity*” and/or “*similarity*” approach
 - *Discontinuity*: segment into regions based on discontinuity (gradient or edge detection)
 - *Similarity*: Merge similar regions (clustering, region growing etc.)
- Topics
 - Simple Segmentation
 - Segmentation by Clustering
 - Superpixel Segmentation

2

2

Goal




3

3

Recap: Otsu's Simple Segmentation

- Distribution of graylevels can be used to determine binary threshold
- Histogram graphs number of pixels in the image with a particular graylevel, as a function of the possible graylevels
 - Find peaks and set threshold between peaks



4

4

Otsu's Method

- “A threshold selection method from graylevel histograms”, IEEE Trans on Sys., Man, and Cyb., Vol 9, No 1, pp 62-66, 1979.
 - Basic idea: threshold is chosen such that the division in the histogram yields the largest reduction in standard deviation of the pixel intensities (black, white)
 - Matlab: `graythresh()`




5

5

Image Segmentation by Clustering

Identify groups of pixels that “go together”

Each “point” is a pixel in color space (3-D: RGB)



- K-Means
- Mean-Shift Clustering

6

6

K-Means

7

7

K-means Clustering

- Each “point” is a 3-D vector of color (RGB)
- Initialization:
 - Choose k cluster center points (how pick k ?)
- Repeat:
 - **Assignment step:**
 - For every point, find its closest center point
 - **Update step:**
 - Update every center point as the mean of its assigned points
- Until:
 - The maximum number of iterations is reached, or
 - No changes during the assignment step, or
 - The average distortion per point drops very little

8

8

K-means: Initialization

- K-means is *extremely sensitive* to initialization
- Bad initialization can lead to
 - Poor convergence speed
 - Poor overall clustering
- How to initialize?
 - Randomly from data
 - Try to find K “spread-out” points (K-means++)
- Try multiple initializations and pick best result
 - Minimize total “distortion/inertia” (sum of distances of points to their cluster centers)

$$J(\mu, r) = \sum_{n=1}^N \sum_{k=1}^K \delta_{nk} \|x_n - \mu_k\|^2$$


Data Cluster center

Whether x_j is assigned to μ_i

9

9

Example



10

10

Mean-Shift

11

11

Mean-Shift segmentation




- Recall Mean-Shift tracking lecture...
- Used here for clustering
 - Unlike K-means, do **not** need initial 'K' number
- Assigns the data points to clusters iteratively by shifting points towards the “local modes”
 - Mode: The highest density of data points in the region, in the context of the Mean-Shift

12

12

Mean-Shift clustering

- **Related to “Kernel Density Estimation” (KDE)**
 - Imagine the data is sampled from a probability distribution
 - Goal: Estimate the underlying distribution (also called the probability density function) for a set of data
 - Place kernel (think “weighting function”) on each point
 - Adding all the individual kernels generates a probability surface (e.g., density function)





13

13

Mean-Shift Clustering

- **Idea:**
 - Make points climb up the hill to the nearest peak on the density surface
 - *Iteratively* shift each point “uphill” until it reaches a peak



14

14

Mean-Shift Algorithm

- Define datapoint x
 - Color only: $[R, G, B]$
 - Spatial and Color: $[x-loc, y-loc, R, G, B]$
- For each datapoint x , find the neighboring points $N(x)$ of x , given Kernel function/window K
- For each x , calculate the **mean shift** $m(x)$:

$$m(x) = \frac{\sum_{x_i \in N(x)} K(x_i - x) x_i}{\sum_{x_i \in N(x)} K(x_i - x)}$$

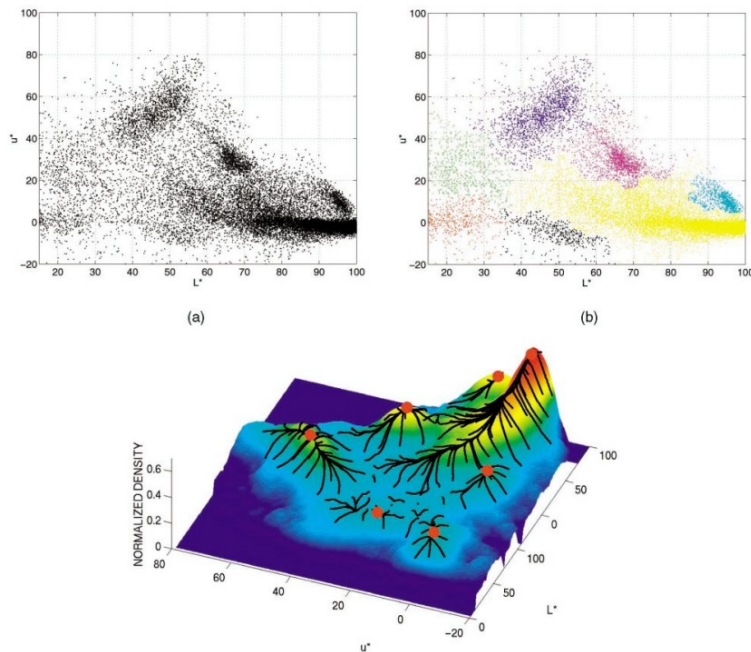
- Then update each x with $x \leftarrow m(x)$
- Repeat n times, or until the points stabilize

More details (Matt Nedrich):

<https://spin.atomicobject.com/2015/05/26/mean-shift-clustering/>

15


15



16

16


RGB Visualization



17

17

Examples



18

18

Superpixels

19

19

Superpixel Segmentation

- “Superpixels” capture local visual redundancy in the image
 - SLIC Superpixel algorithm



20

20

SLIC Superpixel Segmentation

- Generated by clustering pixels based on:
 - Color similarity, and
 - Spatial proximity in the image
- Employ 5-D vector per pixel: $[L, A, B, x, y]$
 - $[L, A, B]$ = pixel color vector in CIELAB color space
 - (L): intensity, (A, B): color
 - “Perceptual color space” with Euclidean distance properties
 - x, y = pixel position (or col, row)

21

21

SLIC Superpixel Algorithm

- Initially choose K = number of desired superpixels
- Divide image into regular “grid” steps S (for the desired K)

Algorithm 1 Efficient superpixel segmentation

- 1: Initialize cluster centers $C_k = [l_k, a_k, b_k, x_k, y_k]^T$ by sampling pixels at regular grid steps S .
- 2: Perturb cluster centers in an $n \times n$ neighborhood, to the lowest gradient position.
- 3: **repeat** (to the pixel with smallest gradient magnitude in 3×3 region)
- 4: **for** each cluster center C_k **do**
- 5: Assign the best matching pixels from a $2S \times 2S$ square neighborhood around the cluster center according to the distance measure (see next few slides)
- 6: **end for**
- 7: Compute new cluster centers and residual error E { $L1$ distance between previous centers and recomputed centers}
- 8: **until** $E \leq \text{threshold}$

22

22

Notation

N	Number of pixels in the input image
K	Number of Superpixels used to segment the input image
N/K	Approximate size of each superpixel
$S = \sqrt{N/K}$	For roughly equally sized superpixels there would be a superpixel centre at every grid interval S

23

23

SLIC Superpixel Segmentation

- Distance function between 2 pixels:

$$D_s = d_{\text{lab}} + \frac{m}{S} d_{x,y}$$

d_{lab} : LAB distance (Euclidean) between the 2 pixels.


$\frac{1}{S} d_{x,y}$: Euclidean spatial distance, normalized by grid interval S.

m : **compactness** control of a super pixel. Larger values make it more compact.

24

24

Superpixels: More Examples



Approx. 300 / 100 superpixel versions

25

25


Template Matching

26

26

Template Matching Intro

- Want to find areas of a search image that are similar to a given template T



27

General Approaches

- Template-Based:
 - Utilize raw template (pixels) and find best matching patches in search image
 - Sum-of-absolute differences (SAD)
 - Sum-of-squared differences (SSD)
 - Normalized cross-correlation (NCC)

28

1) Sum-of-Absolute Differences (SAD)

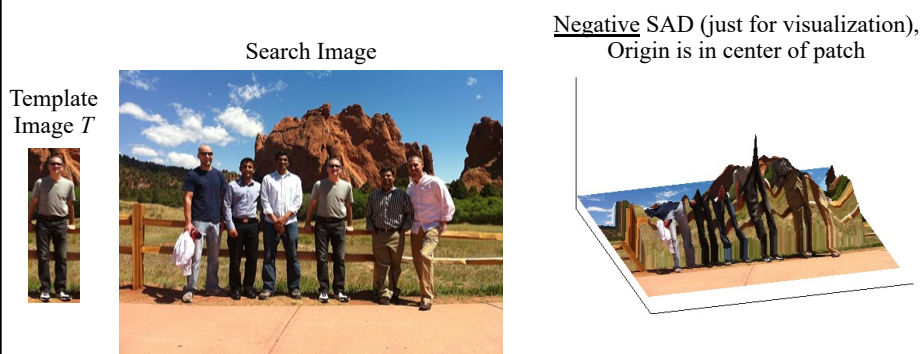
- Compute **absolute differences of pixel intensities** of template T and image patch P extracted from search image (note that P is the same size as template T)

$$SAD(P, T) = \sum_{R, G, B} \sum_{x, y} |P(x, y) - T(x, y)|$$

- Compute SAD for all patch locations within the search image
- Keep patch with minimum SAD or patches with SAD less than given threshold

29



SAD Example



30

SAD Example

k^{th} best matching patch



*Pixel values scaled between 0 and 1

Similar???

31

2) Sum-of-Squared Differences (SSD)

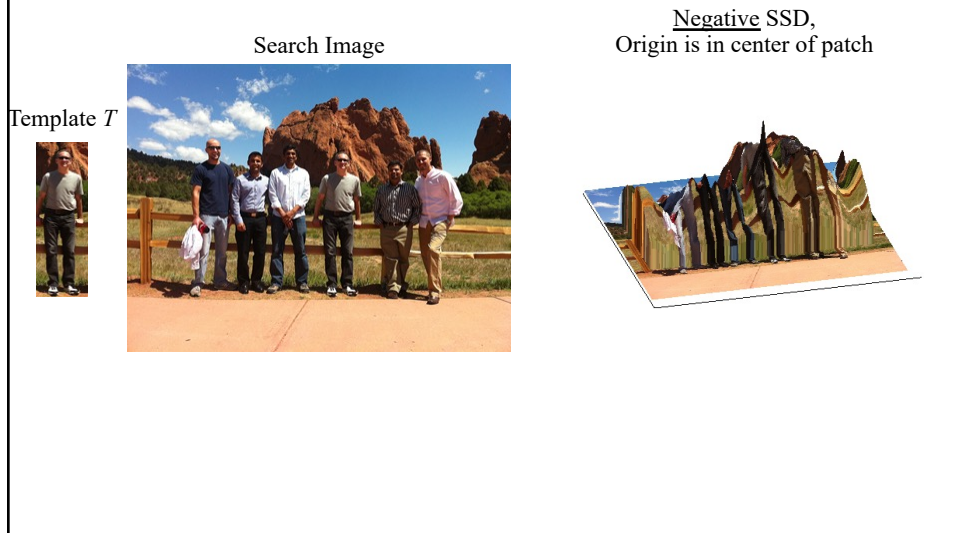
- Similar to SAD, but replace absolute differences with **squared differences**

$$SSD(P, T) = \sum_{R, G, B} \sum_{x, y} (P(x, y) - T(x, y))^2$$

- Compute SSD for all patches within the search image
- Keep patch with minimum SSD

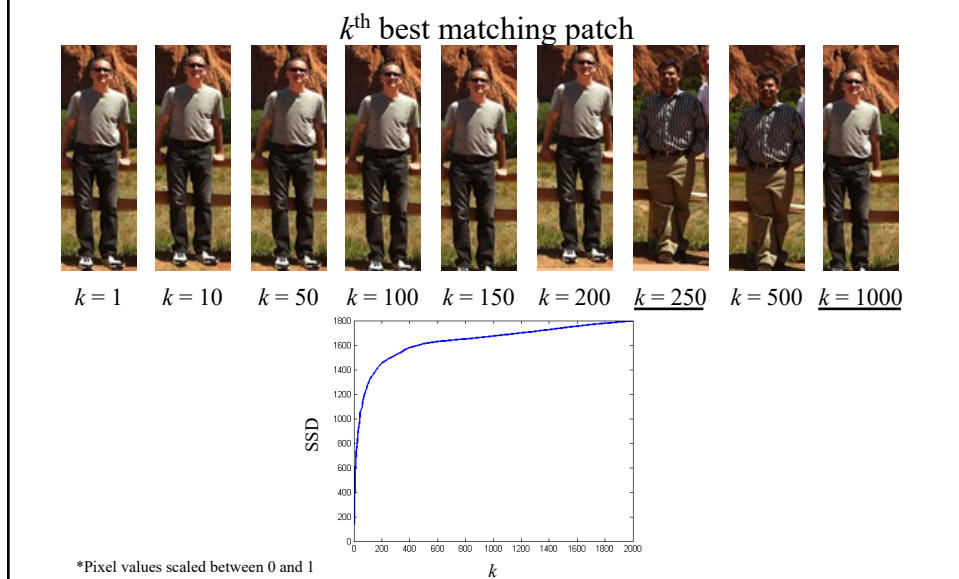
32

SSD Example



33

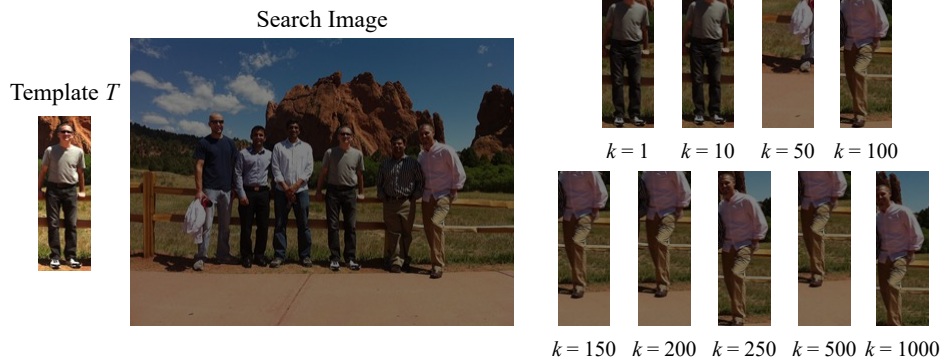
SSD Example



34

Illumination Changes

- SAD and SSD can work well if the template and search image have the same brightness
 - Problem: images can have varying illumination conditions



35

3) Normalized Cross-Correlation (NCC)

- Normalize images to remove variations from illumination conditions

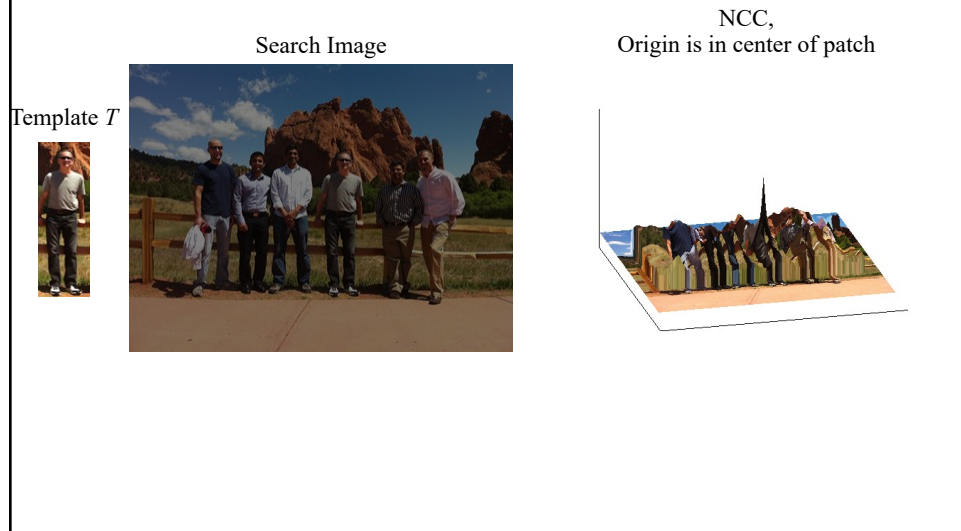
$$NCC(P, T) = \sum_{R, G, B} \frac{1}{n-1} \sum_{x, y} \frac{(P(x, y) - \overset{\substack{\text{Mean of pixel values in patch} \\ \text{(each color computed independently)}}}{\bar{P}}) \cdot (T(x, y) - \overset{\substack{\text{Mean of pixel values in patch} \\ \text{(each color computed independently)}}}{\bar{T}})}{\underset{\substack{\text{Standard deviation of pixel values in patch} \\ \text{(each color computed independently)}}}{\sigma_P \sigma_T}}$$

Note: larger values of NCC better!

The maximum value is 1 when two 1-channel signals are exactly the same: $NCC(a, a)=1$

36



NCC Example



37

NCC Example



k^{th} best matching patch using color images



38

NCC Example


k^{th} best matching patch using grayscale images



39

Handling Differing Scales

- Construct **fixed-sized** template of the smallest size you want to detect
- Scan through image pyramid
 - Detects larger scales of object higher in the pyramid
- Efficient scanning method, with less pixels to examine overall
 - Instead of repeatedly scaling up the template and scanning original full-sized image multiple times



40

Summary

- K-Means
 - Choose cluster centers and label every pixel based on its nearest neighbor
 - Minimize total distortion
 - Sensitive to initialization
- Mean-Shift
 - Iteratively shifts data towards peaks
- Superpixel Segmentation
 - Clustering small pixel regions based on color similarity and proximity
- Template Matching
 - SAD, SSD, NCC

41