

Computer Vision for HCI

Classification Intro

1

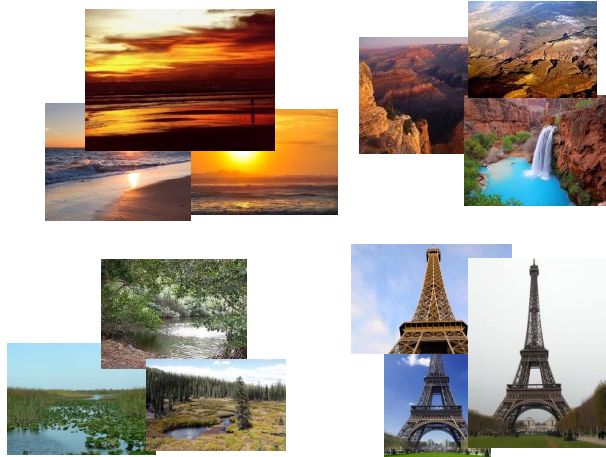
Objective

- To which class does this image belong?

Test Image



Known Classes



2

Supervised Learning

- Given: training data
 - Set of data with corresponding class labels
 - Needs to be representative of entire space of data
- Objective: build a classifier to predict output labels (classes) of data in unseen test set
 - Need to infer a function that separates the data into desirable classes
 - No single algorithm works best on all datasets
 - Need to tune algorithm parameters
 - Feature representation is important

3

Supervised Learning Process

- Split data into training and testing sets
- Determine features to employ
- Select a classifier
- Train the classifier using the training set
- Classify the test set
- Evaluate the classification results

No “data leakage”!

4

Training Classifiers

- Most algorithms have parameters to tune
- Want to avoid over-fitting training data
 - Fitting to noise, not generalizing
- How to tune?
 - Use “validation” data (hypothesized test data)
 - Train classifier on a **subset** of the training data
 - Evaluate the classifier on the **remaining training** data
 - Called the **validation set**
 - Tune the classifier to minimize the error on the validation set

5

Training Classifiers (continued)


- How to tune? (continued)
 - m*-Fold “*Cross-Validation*”
 - Pick/Set classifier options
 - e.g., parameters, model form, training time, or input features
 - Estimate **generalized** classifier performance
 - Randomly divide training set into *m* disjoint sets of equal size
 - Train using (*m*-1) subsets and validate on the remaining subset
 - Repeat *m* times, using different validation set each time
 - Average results
 - Repeat entire process for different classifier options and choose the options which maximize the average results

“Cross-validation is used to estimate the skill of a model”

“The purpose of cross-validation is model checking, not model building”

6

5-Fold Cross-Validation



7

Evaluation

- $\text{Accuracy} = \frac{\text{Number of correct classifications}}{\text{Number of classifications}}$
- Consider the binary classifier situation where we are trying to detect instances of class X within a dataset containing instances of X (positive class) and Y (negative class)
 - True Positive (TP) – Correctly classifying an instance of X as X
 - False Positive (FP) – Incorrectly classifying an instance of Y as X
 - False alarm or Type I error
 - True Negative (TN) – Correctly classifying an instance of Y as Y
 - False Negative (FN) – Incorrectly classifying an instance of X as Y
 - Misdetction or Type II error

8

Evaluation

- **Precision** = $\frac{\text{Number of correctly detected events}}{\text{Number of detected events}} = \frac{TP}{TP + FP}$
- **Recall** = $\frac{\text{Number of correctly detected events}}{\text{True number of events}} = \frac{TP}{TP + FN}$
- **F_β -Measure** = $(1 + \beta^2) \cdot \frac{\text{Precision} \cdot \text{Recall}}{\beta^2 \cdot \text{Precision} + \text{Recall}}$

Common to use $\beta = 1 \rightarrow$ Harmonic mean between precision and recall

9

Classification Approaches


- **k-Nearest Neighbors (kNN)**
- **Decision Trees**

10

k -Nearest Neighbor

- One of the simplest classification strategies
- Algorithm:
 - Compute distance from test sample to labeled training samples
 - Assign test sample the label most common across the first k nearest neighbors from the training data
 - k typically small and odd numbered (no ties)

11



$K=1$ yields X is class o



$K=3$ yields X is class $+$

$K=5$ yields X is class o

12

Example


- Training data (two Gaussians):
 - Class 1 - 667 points sampled from $N\left(\begin{bmatrix} 0 \\ 2.25 \end{bmatrix}, \begin{bmatrix} 5 & 4 \\ 4 & 5 \end{bmatrix}\right)$
 - Class 2 - 333 points sampled from $N\left(\begin{bmatrix} 0 \\ -2.25 \end{bmatrix}, \begin{bmatrix} 5 & 4 \\ 4 & 4 \end{bmatrix}\right)$



13

Example


- Testing dataset consists of 2000 points sampled from same distributions with same priors (2/3 vs. 1/3)



14

Example – k -NN

Classify each testing point using nearest neighbors from training data (known classes)



15

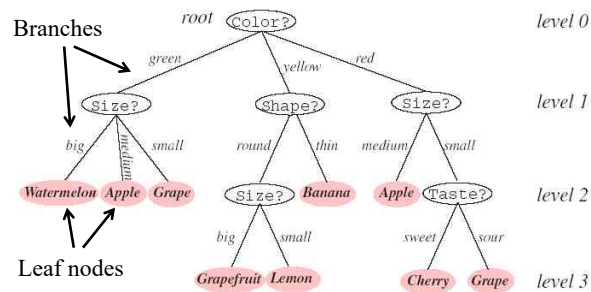
Classification Approaches

- k-Nearest Neighbors (kNN)
- **Decision Trees**

16

Decision Tree

- Classify pattern through sequence of questions
- Easy to interpret



17


CART

- Classification and Regression Trees (CART)
 - General framework for creating decision trees

18

How many splits?

- Every non-binary numeric decision can be represented as combination of binary decisions



19


Which property to test?

- Prefer decisions that lead to simplest tree (Occam's Razor)
 - Want property to split data into "purest" groups possible
 - Use impurity (randomness) measures

Node N

Proportion of patterns at node N that belong to class ω_j

$$\text{Entropy Impurity } i(N) = - \sum_j P(\omega_j) \log_2 P(\omega_j)$$



- Choose decision at node N that decreases impurity the most

Want $\ll i(N)$

$$\text{Maximize } \Delta i(N) = i(N) - [P_L \cdot i(N_L) + (1 - P_L) \cdot i(N_R)]$$

Proportion of patterns that go out to node N_L

Impurity of patterns at node N_L

The diagram shows a node N in an oval, with two arrows pointing down to child nodes N_L and N_R , also in ovals.

20

When to stop?

- Methods to determine when to stop splitting may declare a node a leaf too early
- Alternative: grow tree out entirely (each leaf perfectly pure) and then prune
- Pruning:
 - Work bottom-up
 - Compute the increase in impurity if two child nodes linked to common parent node are eliminated
 - Merge if increase is negligible

21

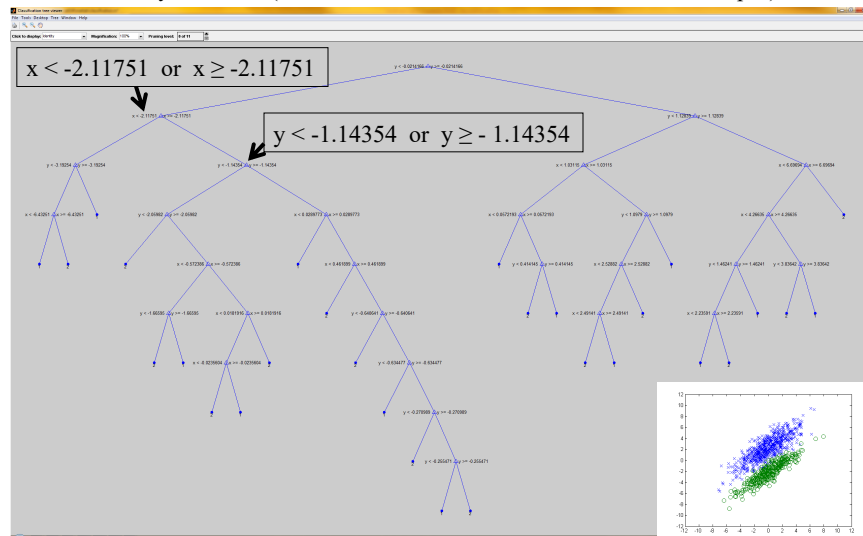
How to assign categories to leaf nodes?

- Simplest approach is to take majority vote of class labels at leaf node
 - Ideally there will be one dominant class
- Potential options when tie occurs:
 - Random assignment
 - Take into account priors
 - Take into account classification risks
 - Cost of misdetections or false alarms of categories

22

Example – Decision Tree (Matlab)

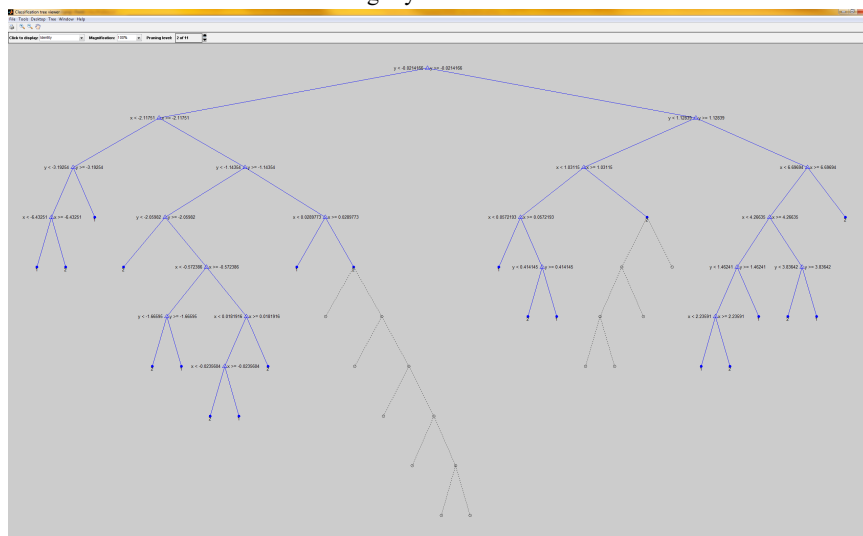
Mostly Full Tree (nodes must have at least 10 observations to be split)



23

Example – Decision Tree

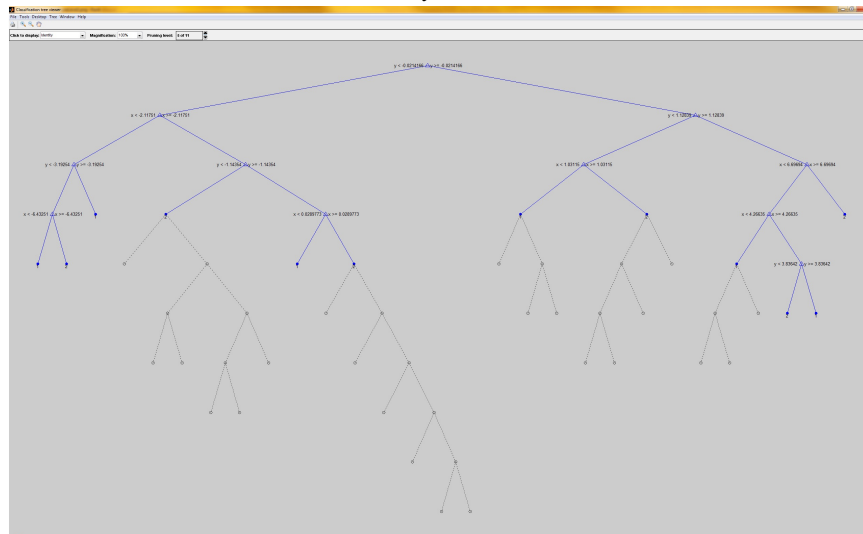
Slightly Pruned Tree



24

Example – Decision Tree

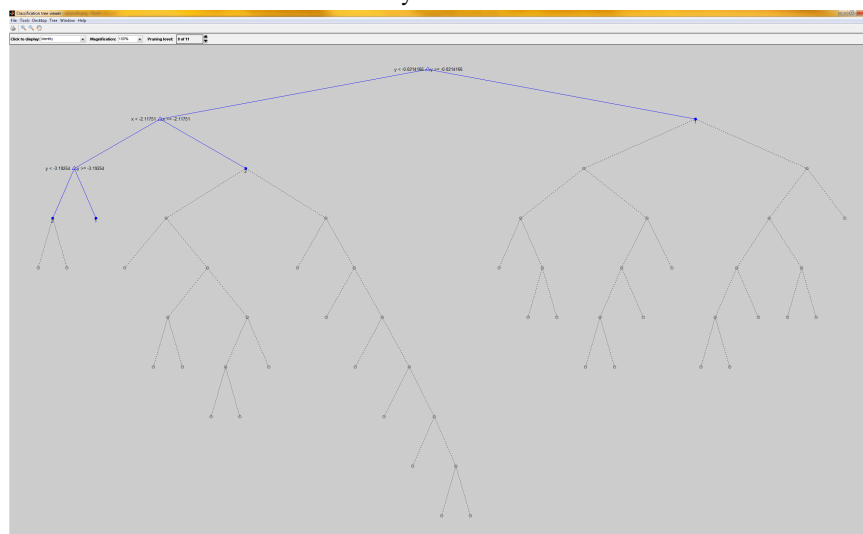
Moderately Pruned Tree



25

Example – Decision Tree

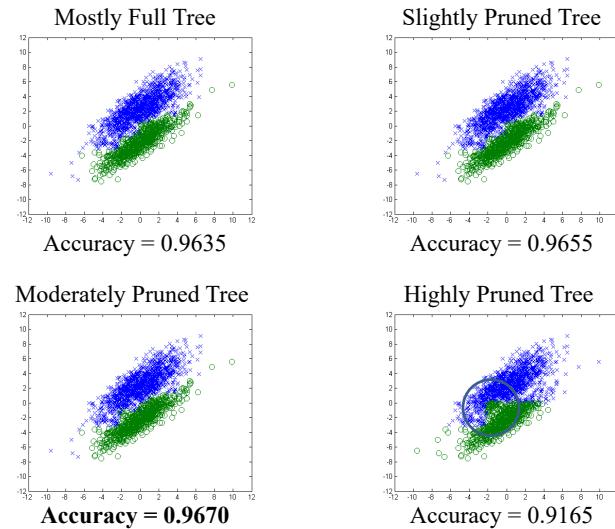
Heavily Pruned Tree



26

Example – Decision Tree

Results from test data



27

Extra: UNsupervised Learning

- Given: ~~training~~ data
 - Set of data WITHOUT corresponding class labels
 - Still needs to be representative of possible space of data
- Primary objective: Group the data into “clusters”
 - Need to measure of closeness/proximity to determine which examples are closer than others
 - Feature representation is still important
- Example: K-means clustering

28


K-means Clustering

- Each “point” is a 3-D vector of color (RGB)
- Initialization:
 - Choose k cluster center points (how pick k ?)
- Repeat:
 - **Assignment step:**
 - For every point, find its closest center point
 - **Update step:**
 - Update every center point as the mean of its assigned points
- Until:
 - The maximum number of iterations is reached, or
 - No changes during the assignment step, or
 - The average distortion per point drops very little

29

29

Previous Segmentation Example



30

30

Summary

- Supervised training concepts
 - Train/Val/Test
 - Cross Validation
- k NN
 - Majority vote of nearest neighbors in training data
- Decision Trees
 - CART